
NEXTXR DIGITAL TWIN

Digital Twin Features

108 features across the platform — what is built, what needs porting, what is not started

Generated 7 September 2026, from the running system

Repository [Tejesh3305/Goalcert_Digital-Twin](#)

Source [docs/DIGITAL-TWIN-FEATURES.md](#)

Status Reflects the deployed implementation, not a design proposal

The target. Workforce Intelligence becomes the main application: an organisation signs in, every person lands in the workspace for their role, and the digital twin, the simulation engine and the agentic AI are capability surfaces reached *through* that workspace at the access level their role allows.

What this document is. Every feature that platform needs, with an honest status against two codebases:

	Meaning
✔ Twin	Built and working in <code>nextxr-ontology-v3</code> (this repo)
◐ Twin partial	Exists here, but narrower than the platform needs
◑ WI only	Built in <code>workforce-intelligence</code> , needs porting or integrating
○ Not built	Exists in neither — genuinely new work

That third status matters more than it looks. A large amount of what follows is already written in `workforce-intelligence/backend/src` and needs *moving*, not inventing — costing integration effort rather than design effort. Treating it as unbuilt would badly overstate what is left.

Surveyed 2026-09-07 against `nextxr-ontology-v3` (27 tables, 188 API routes) and `workforce-intelligence` (15 work tables, 7 roles collapsed to 5 here, 6 role applications, a 9-domain scenario engine).

1. Identity, hierarchy and access

The platform's spine: one account model, a role per person, and every surface rendered from that role's capability list.

#	Feature	Status	Notes
1.1	Organisations, users, sign-in, sessions	✓ Twin	JWT + rotating refresh, session revocation, logout, MFA field
1.2	API keys, scoped per tenant	✓ Twin	Hashed, role-carrying, revocable
1.3	Audit log of account actions	✓ Twin	<code>audit_log</code> , 190 rows live
1.4	Goalcert Hub SSO	✓ Twin	Ticket exchange with replay guard (<code>hub_sso_jti</code>)
1.5	Data role: owner / admin / write / read	✓ Twin	<code>memberships.role</code> , enforced in <code>tenancy.py</code>
1.6	Five-role hierarchy	🕒 Twin	Twin has two personas (supervisor, frontline). The platform needs five — see below
1.7	Capability matrix per role	🕒 Twin	Twin's <code>work/authority.py</code> covers dispatch only. The five roles need ~20 capabilities across content, clearance, evidence, KPI and org
1.8	Team scoping ("whose people?")	✓ Twin	<code>assignable_user_ids</code> — supervisor is authoritative over their own team only
1.9	Frontend renders from the capability list	✓ Twin	<code>GET /work/me</code> returns capabilities; buttons drawn from it
1.10	Approval routing by role	🕒 WI only	<code>APPROVER_FOR</code> routes each approval kind to the right inbox — retargets to Manager / CEO
1.11	Multi-tenant isolation	✓ Twin	<code>org_tenants</code> + per-label <code>(tenantId, id)</code> graph constraints

The five roles

One line of command plus one platform role, which is all the hierarchy this product needs:

Admin • platform: provisioning, connectors, health – not in the line

```

CEO  →  Manager  →  Supervisor  →  Frontline operator
org    site /    one team    the job in front
wide  function                                     of them

```

Role	Owns	Core capabilities
Frontline operator	The job in front of them	act on own assignments, act on work orders, request assist
Supervisor	One team	create assignments (own team), create/close work orders, respond to assist, approve clearance exceptions, read team
Manager	A site or function, several teams	everything a supervisor has but across their teams, plus author/review/publish content, revoke clearances, read and export evidence, read agent logs
CEO	The organisation	set KPI targets, read evidence and readiness org-wide, manage case studies, create assignments
Admin	The platform, not the shift	org management, connector management, platform health

This is a deliberate simplification of what WI ships. WI has seven roles: it splits `lnd` (authoring) and `compliance` (clearance and evidence) into their own roles, and separates `coo` from `super_admin`. Collapsing to five folds **L&D and compliance into Manager**, maps `coo` → **CEO**, and merges `super_admin` into **Admin**.

The capabilities do not disappear — they are redistributed. Porting is therefore a re-keying of WI's **CAPABILITIES** dict rather than a redesign, and the two capabilities that were split across `lnd` and `compliance` (`content.publish` , `clearance.revoke`) simply land on the same role.

Worth knowing before you commit to five. The one thing the merge gives up is *separation of duties*: with Manager holding both `content.publish` and `clearance.revoke` , the person who authors a procedure can also clear people on it and revoke the clearance. In a regulated audit that is usually a finding. It is entirely reasonable to accept for now — the fix later is one extra role, not a redesign, because the capability matrix is data. Flagging it so the choice is made rather than inherited.

The gap that matters here is 1.6/1.7, and everything else hangs off it. The twin's persona axis was built for a two-role dispatch loop. Five roles and a wider capability matrix should be done first; most features below gate on it. The twin's existing two-axis design — `role` for data access, `persona` for job — extends to five personas without structural change.

Also worth keeping: a `role='read'` supervisor currently cannot dispatch, because the auth middleware refuses every POST from a read-only principal before the persona check runs. With five roles that inconsistency gets more expensive, not less.

2. Per-person dashboards

Each role lands in its own workspace. Five roles, five applications — WI already ships six shells, so this is a merge rather than new construction.

#	Application	Role	Status	What it is
2.1	Field	Frontline operator	○ WI + ☒ Twin	Today's jobs, the procedure, the fix. The twin's is the better of the two
2.2	Command	Supervisor	○ WI + ☒ Twin	Dispatch queue, team load, sign-off. Built in the twin
2.3	Manage	Manager	○ WI only	Merge of WI's Studio (authoring, review, publish) and Evidence (clearances, evidence packs, export, agent logs), plus multi-team dispatch
2.4	Exec	CEO	○ WI only	KPIs, targets, readiness, case studies, org-wide rollup
2.5	Control	Admin	○ WI only	Org provisioning, connectors, platform health
2.6	Role-based routing, lazy-loaded per app	○ WI only	PersonaRouter — a session downloads only its own application	
2.7	"No workspace for this role" handling	○ WI + ☒ Twin	Both treat an unassigned role as a provisioning bug, not a crash	
2.8	Operator dashboard with charts and live 3-D	☒ Twin	KPI tiles, completion bars, severity mix, score trend, XP, and the job's twin rendered beside it	
2.9	Supervisor team-load view	☒ Twin	Per-operator load and throughput, org-wide with assignable flags	
2.10	Manager multi-team view	○ Not built	The twin's team view is single-supervisor scoped; a Manager spans several teams	

Reading of this section: the two operational workspaces are built, and the twin's are ahead of WI's — charts, a live 3-D asset panel, click-through to the plant. The three management workspaces exist only as WI shells. The work is porting the twin's dashboard standard into them, not building dashboards from nothing.

On merging Studio and Evidence into one Manage app: they were separate in WI because they served separate roles. Under five roles one person does both, and two applications for one person is two doors to the same room. Manage should be one workspace with sections, not a tab bar joining two former apps.

3. Work: dispatch and execution

#	Feature	Status	Notes
3.1	Teams and team membership	✓ Twin	
3.2	Assign work to a person, team-scoped	✓ Twin	Enforced at service and route level; tested
3.3	Raise work by hand (not from a detection)	✓ Twin	POST <code>/work/tasks</code> , exposed in Dispatch
3.4	Task lifecycle: open → assigned → in progress → resolved → closed	✓ Twin	Legal transitions enforced in one module
3.5	Blocked state with a reason	✓ Twin	
3.6	Append-only transition history	✓ Twin	<code>task_events</code> — answers "who assigned this, when"
3.7	Operator inbox, own work only	✓ Twin	No user id parameter; server reads the session
3.8	Supervisor sign-off closes the loop	✓ Twin	Closing resolves the Finding back in the graph
3.9	Assignment and WorkOrder as separate objects	🕒 Twin	Twin merges both into <code>tasks</code> . WI models them separately — an assignment is <i>what you owe</i> , a work order is <i>the authorised job</i>
3.10	Assist: request help, respond, close	🕒 WI only	<code>AssistSession</code> — an operator escalates to a supervisor mid-job
3.11	Notifications	🕒 Not built	WI has the <code>Notification</code> table and endpoints; no delivery channel exists in either (no email/push/SMS)
3.12	Work rules: which findings become work	✓ Twin	Data-driven; default critical-only, dedup by finding node
3.13	Auto-generated work order content	✓ Twin	AI-written, on explicit request

4. Competency, clearance and evidence

The largest genuinely missing area in this repo. Nothing here exists in the twin; all of it exists in WI.

#	Feature	Status	Notes
4.1	Clearance: a person is cleared to do a procedure on an asset class	○ WI only	<code>Clearance</code> model
4.2	Evidence links backing a clearance	○ WI only	<code>EvidenceLink</code> — what proves competency
4.3	Clearance revocation	○ WI only	Manager capability under the five-role model
4.4	Clearance exception approval	○ WI only	Routed to Supervisor or Manager
4.5	Evidence pack export	○ WI only	The audit deliverable
4.6	Clearance gates dispatch	○ Not built	Neither codebase refuses an assignment because the assignee is not cleared. This is the point of the feature and it is unbuilt in both
4.7	Expiry and recertification	○ Not built	No scheduled re-validation anywhere
4.8	Agent action log for audit	● Twin	Twin meters agent spend and records <code>task_events</code> ; WI has <code>agentlog.read</code> as a capability. Neither writes a reviewable per-decision agent log

4.6 is the one to notice. Clearance data without a gate is a record nobody is forced to consult. The moment `assign` refuses an uncleared operator, the whole competency model starts paying for itself — and that check belongs next to `assignable_user_ids` , which already exists in the twin.

5. Learning, content and the improvement loop

#	Feature	Status	Notes
5.1	Procedure catalogue	🕒 Twin	Twin has 980 lines of hand-authored procedures in code ; WI has a Procedure table. Content in code cannot be authored by an L&D user
5.2	Graded decision drills with server-side scoring	✅ Twin	Answer key never leaves the server; tested
5.3	Practice vs assessed runs	✅ Twin	
5.4	Hints, wrong-step penalties, debrief	✅ Twin	
5.5	XP, levels and an auditable ledger	✅ Twin	Twin-only — WI has no XP model
5.6	Training library with per-item standing	✅ Twin	Passed / in progress / attempts / last score; resume a half-finished run
5.7	"Repair with AI" reusable trainer	✅ Twin	Component reachable from the job, the library and the twin
5.8	Candidate content pipeline: draft → review → publish	🕒 WI only	<code>CandidateContent</code> + <code>content.publish</code> / <code>content.review</code>
5.9	Authoring UI for procedures	🕒 WI only	WI's Studio app; becomes a section of Manage
5.10	Loop events: what the fleet learned	🕒 WI only	<code>LoopEvent</code> — feeding outcomes back into content
5.11	Curriculum / learning paths	🕒 Not built	No ordered multi-procedure programme in either
5.12	Skills matrix per person	🕒 Not built	Implied by clearances; no aggregate view

6. Scenario builder and simulation engine

You asked specifically for the scenario builder. It exists — in WI, not here.

#	Feature	Status	Notes
6.1	Scenario builder / authoring studio	⊖ WI only	<code>scenario_engine/api/studio.py</code> — authors a scenario from a natural-language description , browses domains and fault types, runs it
6.2	Scenario model: objectives, steps, decision gates, target selectors	⊖ WI only	<code>engine/scenario.py</code>
6.3	Monte Carlo runs	⊖ WI only	<code>engine/monte_carlo.py</code> — distribution of outcomes, not one path
6.4	World model: actors, resources, environment	⊖ WI only	<code>engine/world.py</code> , <code>environment.py</code>
6.5	Workflows and conditions	⊖ WI only	<code>engine/workflows.py</code> , <code>conditions.py</code>
6.6	KPI evaluation per run	⊖ WI only	<code>engine/kpis.py</code>
6.7	Nine domain plug-ins	⊖ WI only	aerospace, defence, edm, ev, facility, fleet, hospital, railway, solar
6.8	Run manager, run history, tripwires	⊖ WI only	<code>api/runs.py</code> , <code>api/tripwire.py</code>
6.9	Scenario reports	⊖ WI only	<code>reports/generator.py</code>
6.10	Guided mode	⊖ WI + ✅ Twin	Both have a guided path; the twin's is the graded operator drill
6.11	Scenario engine reads live twin state	⊖ WI only	<code>services/twin_client.py</code> — the seam that makes scenarios run against a real plant
6.12	Physics simulation of a machine	✅ Twin	Different thing, and the twin is ahead: <code>dynamics/</code> integrates wear, faults persist, one lease-owner per tenant ticks it

Two different simulators, and they are complementary rather than duplicated. The twin's `dynamics/` engine answers *what the machine does* — a physical integrator where wear accumulates and an injected fault persists. WI's `scenario_engine` answers *what the organisation does about it* — actors, resources, decision gates and a Monte Carlo spread of outcomes. The platform needs both, wired together through `twin_client.py`. Neither replaces the other, and the deliberate decision not to port WI's engine into the twin still stands; what changes is that WI becomes the host, so the engine is already in the right place.

7. The digital twin as a capability surface

This is the area the twin is furthest ahead on, and none of it exists in WI's older embedded copy at the same standard.

#	Feature	Status	Notes
7.1	Ontology-driven graph: 11 node categories, SHACL shapes	✓ Twin	
7.2	Entity CRUD with a hash-chained change log	✓ Twin	Tamper-evident
7.3	Live telemetry ingest (device tokens, bulk, idempotent)	✓ Twin	Composite PK makes replay safe
7.4	Protocol connectors + SunSpec discovery	✓ Twin	14 endpoints
7.5	Historian: history, trends, signal discovery	✓ Twin	34,347 measurements live
7.6	Behaviour registry raising Findings	✓ Twin	
7.7	Live 3-D scenes per domain	✓ Twin	Including a turbine model and a hospital floor plan
7.8	Photo → 3-D reconstruction	✓ Twin	Via RunPod GPU
7.9	Twin builder from natural language	✓ Twin	Agent-driven
7.10	Per-tenant runtime ownership via Redis lease	✓ Twin	Prevents duplicate writes at 2+ tasks
7.11	SSE event bus per tenant	✓ Twin	
7.12	Prediction / RUL and cascade analysis	✓ Twin	
7.13	Twin access scoped by workforce role	○ Not built	Today twin access is <code>role</code> + <code>org_tenants</code> . Nothing expresses "a frontline operator may see only the assets their team owns"
7.14	Asset ↔ twin binding from the work side	○ WI only	WI's <code>POST /work/assets/bind-twins</code> ; the twin has no equivalent
7.15	Graph is populated	○ Blocked	3 nodes against 16 registered twins. Schema and constraints are provisioned; content is not. Everything above is demonstrable only once this is seeded

8. Agentic AI

#	Feature	Status	Notes
8.1	16 embedded copilot agents	✓ Twin	Narration, diagnosis, analysis, cascade, work order, procurement, incident report, procedure, troubleshoot, chat
8.2	Deterministic stub for every agent	✓ Twin	Keyless operation never silently fakes a real answer
8.3	Spend ledger, response memo, in-flight coalescing	✓ Twin	Per-model token and cost accounting
8.4	Per-tenant budget cap	🕒 Twin	Implemented but disabled unless <code>NXR_COPILOT_BUDGET_USD</code> is set
8.5	Long-running agent sessions	✓ Twin	26 endpoints: twin builder, bundle author, plugin, accelerator, ops
8.6	Scenario analyst agent	🕒 WI only	<code>embedded_agents/scenario/analyst.py</code>
8.7	Scenario authoring agent	🕒 WI only	<code>embedded_agents/scenario/authoring.py</code> — powers the builder
8.8	Per-role agent access levels	🕒 Not built	Every signed-in user can call every agent. No role gates which agents run, and no per-role budget
8.9	Reviewable agent decision log	🕒 Not built	Spend is metered; individual decisions are not recorded for audit

9. Analytics, KPIs and readiness

#	Feature	Status	Notes
9.1	Operator's own stats: counts, series, score, streak	✓ Twin	<code>GET /work/stats/me</code>
9.2	Supervisor team throughput	✓ Twin	<code>GET /work/team-stats</code>
9.3	XP ledger and standing	✓ Twin	
9.4	KPI points and targets	○ WI only	<code>KpiPoint</code> , <code>kpi.set_target</code>
9.5	Readiness snapshots	○ WI only	<code>ReadinessSnapshot</code> — is the workforce ready for the work
9.6	Org-wide executive view	○ WI only	Exec app
9.7	Case study management	○ WI only	<code>casestudy.manage</code>
9.8	Cross-tenant fleet analytics	○ Not built	Neither aggregates across organisations

10. Platform, deployment and integration

#	Feature	Status	Notes
10.1	Migration ledger and CLI	✓ Twin	27 tables, head <code>0006_work_graph</code>
10.2	SQLite (dev) / MySQL (prod) from one DDL	✓ Twin	
10.3	Multi-task safe with required-flag guards	✓ Twin	<code>NXR_REQUIRE_DB/S3/REDIS</code>
10.4	ECS Fargate + RDS + ElastiCache + S3 + Neo4j on EC2	✓ Twin	Documented in <code>AWS-ARCHITECTURE.md</code>
10.5	Health / readiness split for the load balancer	✓ Twin	
10.6	Federated frontend (twin mounts inside a host shell)	✓ Twin	Module federation — this is what makes WI-as-host feasible
10.7	Two codebases with two data layers	○ Blocked	Twin uses raw SQL over <code>db/core.py</code> ; WI uses SQLAlchemy. One has to win, and it is the single biggest structural decision ahead
10.8	Org provisioning / connector management UI	○ WI only	Control app
10.9	CI, tests	✓ Twin	528 tests passing
10.10	Notification delivery (email / push / SMS)	○ Not built	No transport in either codebase

11. Scorecard

Counted from the tables above, not estimated:

Status	Count	Share
✓ Built in the twin	55	51%
○ Built in WI, needs porting or integrating	34	31%
● Partial in the twin	6	6%
○ Not built anywhere	13	12%
Total	108	

Seven features in eight already exist in one codebase or the other. The dominant cost is not writing features — it is choosing a host, reconciling two data layers, and widening the role model so five roles can reach capabilities that already work for two.

The thirteen genuinely unbuilt

#	Feature	Note
2.10	Manager multi-team view	The twin's team view is scoped to one supervisor
3.11	Notifications	Model exists in WI; nothing delivers
4.6	Clearance gates dispatch	The check that makes the competency model matter
4.7	Clearance expiry and recertification	
5.11	Curriculum / learning paths	
5.12	Skills matrix per person	
7.13	Twin access scoped by workforce role	
7.15	Graph population	Blocked on data, not code
8.8	Per-role agent access levels and budgets	
8.9	Reviewable agent decision log	
9.8	Cross-tenant fleet analytics	
10.7	Unified data layer	A decision, then a migration
10.10	Notification delivery transport	No email/push/SMS in either codebase

Note what is *not* on this list: the five-role model (1.6/1.7) and the Assignment/WorkOrder split (3.9) are **partial**, not missing — the structures exist and need widening. That distinction is the difference between a fortnight and a quarter.

The six partials

#	Feature	What "partial" means here
1.6	Five-role hierarchy	Two personas exist; the axis extends to five without structural change
1.7	Capability matrix	Covers dispatch; needs content, clearance, evidence, KPI
3.9	Assignment vs WorkOrder	Merged into tasks ; WI models them separately
4.8	Agent action log	Spend is metered; individual decisions are not recorded
5.1	Procedure catalogue	980 lines in code — an L&D user cannot author it
8.4	Per-tenant budget cap	Implemented but disabled unless the env var is set

12. Suggested order

Sequenced by what unblocks the most, not by size.

First — decide the host and the data layer (10.7). WI as the application shell with the twin federated in is the stated direction and is already technically supported. The data layer is the real fork: the twin's raw-SQL store carries the migration ledger, MySQL/SQLite duality and 528 tests; WI's SQLAlchemy models carry clearance, evidence, approvals and KPIs. Porting WI's 15 work models onto the twin's `db/` layer is the smaller and better-tested direction, but it is a decision to make deliberately rather than by drift.

Second — widen the role model (1.6, 1.7). Five roles and a full capability matrix. Everything in sections 4, 5 and 9 gates on this, and the twin's existing two-axis design (`role` for data, `persona` for job) extends to it cleanly.

Third — seed the graph (7.15). Nothing in section 7 demonstrates against 3 nodes, and this is a data task, not development.

Fourth — port the three management applications (2.3–2.5) onto the twin's dashboard standard, which is currently the better of the two.

Fifth — bring the scenario engine and builder across (section 6) with `twin_client.py` pointed at this repo's live twin API.

Then the genuinely new work, led by clearance gating (4.6), because it is small, it is the point of the competency model, and it belongs next to `assignable_user_ids`, which is already written and tested.